

```
In [20]: import warnings, requests, zipfile, io
warnings.simplefilter('ignore')
import pandas as pd
from scipy.io import arff
import boto3
```

```
In [21]: f_zip = 'http://archive.ics.uci.edu/ml/machine-learning-databases/00212/vertebra
r = requests.get(f_zip, stream=True)
Vertebral_zip = zipfile.ZipFile(io.BytesIO(r.content))
Vertebral_zip.extractall()
```

```
In [22]: data = arff.loadarff('column_2C_weka.arff')
df = pd.DataFrame(data[0])
```

```
In [23]: class_mapper = {b'Abnormal':1,b'Normal':0}
df['class']=df['class'].replace(class_mapper)
```

```
In [24]: df.shape
```

```
Out[24]: (310, 7)
```

```
In [25]: df.columns
```

```
Out[25]: Index(['pelvic_incidence', 'pelvic_tilt', 'lumbar_lordosis_angle',
               'sacral_slope', 'pelvic_radius', 'degree_spondylolisthesis', 'class'],
              dtype='object')
```

```
In [26]: df.columns
```

```
Out[26]: Index(['pelvic_incidence', 'pelvic_tilt', 'lumbar_lordosis_angle',
               'sacral_slope', 'pelvic_radius', 'degree_spondylolisthesis', 'class'],
              dtype='object')
```

```
In [27]: cols = df.columns.tolist()
cols = [cols[-1]] + cols[:-1]
df = df[cols]
```

```
print(df.columns)
```

```
Index(['class', 'pelvic_incidence', 'pelvic_tilt', 'lumbar_lordosis_angle',
      'sacral_slope', 'pelvic_radius', 'degree_spondylolisthesis'],
      dtype='object')
```

```
In [28]: from sklearn.model_selection import train_test_split
train, test_and_validate = train_test_split(df, test_size=0.2, random_state=42,
```

```
In [29]: test, validate = train_test_split(test_and_validate, test_size=0.5, random_state
```

```
In [30]: print(train.shape)
print(test.shape)
print(validate.shape)
```

```
(248, 7)
```

```
(31, 7)
```

```
(31, 7)
```

```
In [31]: print(train['class'].value_counts())
print(test['class'].value_counts())
print(validate['class'].value_counts())
```

```
class
1    168
0     80
Name: count, dtype: int64
class
1     21
0     10
Name: count, dtype: int64
class
1     21
0     10
Name: count, dtype: int64
```

```
In [32]: bucket='c169682a4380819111150392t1w164112440464-labbucket-divyesujtnis'
```

```
prefix='lab3'
```

```
train_file='vertebral_train.csv'
test_file='vertebral_test.csv'
validate_file='vertebral_validate.csv'
```

```
import os
```

```
s3_resource = boto3.Session().resource('s3')
def upload_s3_csv(filename, folder, dataframe):
    csv_buffer = io.StringIO()
    dataframe.to_csv(csv_buffer, header=False, index=False)
    s3_resource.Bucket(bucket).Object(os.path.join(prefix, folder, filename)).put
```

```
INFO:botocore.credentials:Found credentials from IAM Role: BaseNotebookInstance
Ec2InstanceRole
```

```
In [33]: upload_s3_csv(train_file, 'train', train)
upload_s3_csv(test_file, 'test', test)
upload_s3_csv(validate_file, 'validate', validate)
```

```
In [34]: import boto3
from sagemaker.image_uris import retrieve
container = retrieve('xgboost', boto3.Session().region_name, '1.0-1')
```

```
INFO:sagemaker.image_uris:Defaulting to only available Python version: py3
INFO:sagemaker.image_uris:Defaulting to only supported image scope: cpu.
```

```
In [35]: hyperparams={"num_round": "42",
                      "eval_metric": "auc",
                      "objective": "binary:logistic"}
```

```
In [36]: import sagemaker
s3_output_location="s3://{}/{}/output/".format(bucket, prefix)
xgb_model=sagemaker.estimator.Estimator(container,
                                         sagemaker.get_execution_role(),
                                         instance_count=1,
                                         instance_type='ml.m4.xlarge',
                                         output_path=s3_output_location,
```

```
hyperparameters=hyperparams,  
sagemaker_session=sagemaker.Session())
```

```
In [37]: train_channel = sagemaker.inputs.TrainingInput(  
        "s3://{}/{}/train/".format(bucket,prefix,train_file),  
        content_type='text/csv')  
  
        validate_channel = sagemaker.inputs.TrainingInput(  
        "s3://{}/{}/validate/".format(bucket,prefix,validate_file),  
        content_type='text/csv')  
  
        data_channels = {'train': train_channel, 'validation': validate_channel}
```

```
In [38]: xgb_model.fit(inputs=data_channels, logs=False)
```

INFO:sagemaker.telemetry.telemetry\_logging:SageMaker Python SDK will collect telemetry to help us better understand our user's needs, diagnose issues, and deliver additional features.

To opt out of telemetry, please disable via TelemetryOptOut parameter in SDK defaults config. For more information, refer to <https://sagemaker.readthedocs.io/en/stable/overview.html#configuring-and-using-defaults-with-the-sagemaker-python-sdk>.

INFO:sagemaker:Creating training-job with name: sagemaker-xgboost-2025-08-08-06-16-22-602

2025-08-08 06:16:22 Starting - Starting the training job...

2025-08-08 06:16:46 Starting - Preparing the instances for training....

2025-08-08 06:17:12 Downloading - Downloading input data.....

2025-08-08 06:17:42 Downloading - Downloading the training image.....

2025-08-08 06:18:32 Training - Training image download completed. Training in progress....

2025-08-08 06:18:53 Uploading - Uploading generated training model..

2025-08-08 06:19:06 Completed - Training job completed

```
In [ ]:
```